

8.23-9.5 周报

赵晓辉

2021 年 9 月 5 日

1 写在前面

这个双周主要完成了 cs231n 课程的 Lecture 1 至 Lecture 6, 主要内容包括距离函数、KNN、SVM、损失函数及优化、BP 算法、CNN 架构、非线性激活函数以及神经网络的参数优化等, 并完成 cs231n assignment1。课程概要笔记及 assignment 将在<https://github.com/zxh991103/cs231NOTE>持续跟踪。此外使用 torch, 初步学习和实现了 GAT 算法。

2 Lec 1-6 课程概要

2.1 距离函数

L_1 Distance

$$d_1(I_1, I_2) = \sum_p |I_1^p - I_2^p|$$

L_2 Distance

$$d_2(I_1, I_2) = \sqrt{\sum_p (I_1^p - I_2^p)^2}$$

2.2 KNN

计算测试样本与所有训练集样本之间的距离值, 并根据 K 值投票选举出最相似的标签。

2.3 SVM

计算能够划分训练集样本且距离最大的超平面。

$$w \cdot x + b = 0$$

2.4 损失函数

损失函数评估模型预测值与模型真实值之间的差异性，我们要将其最小化。对于给定的训练集 $(x_i, y_i)_{i=1}^N$ ，我们有损失函数：

$$L = \frac{1}{N} \sum_i L_i(f(x_i, W), y_i)$$

对于 Multi-SVM，我们有损失函数，即 hinge loss：

$$s = f(x_i, W)$$

$$L_i = \sum_{j \neq y_i} \max(0, s_j - s_{y_i} + 1)$$

对于 softmax loss：

$$L_i = -\log\left(\frac{e^{s_{y_i}}}{\sum_j e^{s_j}}\right)$$

2.5 正则化

根据奥卡姆剃刀原则，模型越简单越符合实际，所以我们将正则惩罚项加在损失函数上。

$$L = \frac{1}{N} \sum_i L_i(f(x_i, W), y_i) + R(W)$$

L1

$$R(W) = \sum_k \sum_l |W_{k,l}|$$

L2

$$R(W) = \sum_k \sum_l W_{k,l}^2$$

Elastic

$$R(W) = \sum_k \sum_l \beta W_{k,l}^2 + |W_{k,l}|$$

2.6 BP 算法

链式法则：

$$\frac{\partial f}{\partial y} = \frac{\partial f}{\partial q} \frac{\partial q}{\partial y}$$

故在计算损失函数对于参数的梯度值时，我们应当将本地梯度值与上游回传梯度值相乘。

此时，我们也可以发现 Relu 函数，即 max gate 中只有前向传播计算中的正值能影响下游。

A vectorized example:

$$f(x, W) = \|Wx\|^2 = \sum_{i=1}^n (Wx)_i^2$$

$$q = W \cdot x$$

$$\nabla_W f = 2q \cdot x^T$$

2.7 NN

假设我们将神经网络的计算图表示为:

$$f = \text{Softmax}(W_2 \text{Relu}(W_1 x))$$

神经元 A 拥有 W_1 , 能够具有识别出来 100 种特征的功能, 比如识别出马的左脸或者右脸、车头或者车尾。而神经元 B 拥有 W_2 , 其功能就在于将马的左脸或右脸合并为马的特征, 将车头或车尾合并成车的特征, 从而进行识别。

2.8 CNN

卷积层:

假设有 $32 \times 32 \times 3$ 的图片, 卷积核 $w \ 5 \times 5 \times 3$, 以及偏置 b , 卷积后我们获得 $28 \times 28 \times 1$ 的矩阵, 其中 1 是卷积核的数量。卷积公式为 (每位相乘再求和):

$$f[x, y] * g[x, y] = \sum_{n_1=-\infty}^{\infty} \sum_{n_2=-\infty}^{\infty} f[x, y] * g[x - n_1, y - n_2]$$

步长 (stride):

假设我们有 7×7 的输入, 3×3 的卷积核, 2 的步长, 最后的输出为 3×3 。

此时 $\text{outputsize} = \frac{(N-F)}{\text{stride}} + 1$

填充 (Pad):

图像四周补充 0, 来防止在深层卷积时张量过小。此时, $\text{outputsize} = \frac{(N-F+2P)}{\text{stride}} + 1$ 。

Example:

input volume $32 \times 32 \times 3$, 10 5×5 filters (include 3 depth), stride 1, pad 2, we have 760 parameters ($10 * (5 * 5 * 3 + 1 \text{ bias}) = 760$)

pooling layer: 相当于下采样。

maxpooling:

一般, 每一个池化 filter 具有和步长相同的大小以避免 overlap.

例如,

1	1	2	4
5	6	7	8
3	2	1	0
1	2	3	4

我们使用 2*2 的 filter 和 2 的 stride,maxpooling 后变为:

6	8
3	4

2.9 激活函数

若全部线性连接则等同于一个线性连接, 所以网络中需要非线性的激活函数变换。

sigmod

$$\sigma(x) = \frac{1}{1 + e^{-x}}$$

tanh

$$\tanh(x)$$

Relu

$$\max(0, x)$$

LeakyRelu

$$\max(0.1x, x)$$

Maxout

$$\max(w_1^T x + b_1, w_2^T x + b_2)$$

Elu

$$\begin{cases} x & x \geq 0 \\ \alpha(e^x - 1) & x < 0 \end{cases}$$

SELU

$$f(x) = \begin{cases} \lambda x & \text{if } x > 0 \\ \lambda \alpha(e^x - 1) & \text{otherwise} \end{cases}$$

softmax 的问题:

- x 过大或过小，本地梯度接近 0，使得与上游梯度乘积也接近于 0，更新缓慢。
- x 的值恒正或恒负，本地梯度总是大于 0 的，造成 w 的移动时锯齿状的，接近最优值放缓。

relu 的问题:

- 若 $w \cdot x + b$ 总是负的，则本地梯度为 0，造成参数不更新。

2.10 数据处理

```
1 # ZERO-CENTER
2 X -= np.mean(X,axis = 0)
3 # normalize
4 X /= np.std(X,axis = 0)
```

Listing 1: normalization

2.11 参数初始化

Naive: 为参数初始化小随机数。但是随着网络深度的增加，本地梯度与上游梯度相乘之后接近零，学习十分缓慢。

Xavier:

```
1 W = np.random.randn(dim_in,dim_out)/np.sqrt(dim_in)
```

Listing 2: Xavier

原因:

we want $\text{Var}(y) = \text{Var}(x_i)$, and we have

$$y = \sum_{i=1}^{D_{in}} x_i w_i$$

and we assume that every x has same var. so we have

$$\text{var}(y) = D_{in} \times \text{var}(x) \times \text{var}(w_i)$$

and obviously initial $w_i \sim N(0,1)$, we make $\frac{w_i}{\sqrt{D_{in}}}$ to achieve the var is $\frac{1}{D_{in}}$

Kaiming/MSRA:

```
1 W = np.random.randn(dim_in,dim_out)*np.sqrt(2/dim_in)
```

Listing 3: MSRA

2.12 Batch Normalization

we have the input x like $N \times D$

$$\hat{x}^{(k)} = \frac{x^{(k)} - E[x^{(k)}]}{\sqrt{\text{Var}[x^{(k)}]}}$$

so that we have:

$$\mu_j = \frac{1}{N} \sum_{i=1}^N x_{i,j}$$

$$\sigma_j^2 = \frac{1}{N} \sum_{i=1}^N (x_{i,j} - \mu_j)^2$$

$$\hat{x}_{i,j} = \frac{x_{i,j} - \mu_j}{\sqrt{\sigma_j^2 + \epsilon}}$$

the net is supposed to learn $\gamma \in R^D$ and $\beta \in R^D$

$$y_{i,j} = \gamma_j \hat{x}_{i,j} + \beta_j$$

BN 层经常用于全连接或卷积层后。

2.13 Norm For Conv.

batch norm , 对于每一个 batch 中的每个 channel 取平均得到 $(1 \times 1 \times C)$

layer norm , 对于所有的 batch 我们取一个平均的图片 $(H \times W \times C)$

instance norm , 对于所有的 batch 我们将取一个平均的单通道图片 $(H \times W \times 1)$

group norm , 对于所有 batch, 我们将 channel 分为 k 个组, 在每个组上取平均值得到 $(H \times W \times k)$

3 Assignment1

3.1 KNN

双循环实现距离矩阵:

```
1 num_test = X.shape[0]
2 num_train = self.X_train.shape[0]
3 dists = np.zeros((num_test, num_train))
4 for i in range(num_test):
5     for j in range(num_train):
6         t1 = X[i]
```

```

7         t2 = self.X_train[j]
8         t = t1 - t2
9         t = t*t
10        t = t.sum()
11        t = t**0.5
12        dists[i][j] = t

```

Listing 4: KNN 双循环实现距离矩阵

单循环实现距离矩阵：

```

1    num_test = X.shape[0]
2    num_train = self.X_train.shape[0]
3    dists = np.zeros((num_test, num_train))
4    for i in range(num_test):
5        dists[i,:] = np.sum((X[i,:]-self.X_train)**2,axis = 1)

```

Listing 5: KNN 单循环实现距离矩阵

向量操作实现距离矩阵（矩阵下的完全平方公式）：

```

1    M = np.dot(X, self.X_train.T)
2    nrow=M.shape[0]
3    ncol=M.shape[1]
4    te = np.diag(np.dot(X,X.T))
5    tr = np.diag(np.dot(self.X_train,self.X_train.T))
6    te= np.reshape(np.repeat(te,ncol),M.shape)
7    tr = np.reshape(np.repeat(tr, nrow), M.T.shape)
8    sq=-2 * M +te+tr.T
9    dists = np.sqrt(sq)

```

Listing 6: KNN 向量操作实现距离矩阵

KNN 预测：

```

1    num_test = dists.shape[0]
2    y_pred = np.zeros(num_test)
3    labelnum = len(np.unique(self.y_train))
4    for i in range(num_test):
5        # A list of length k storing the labels of the k nearest neighbors
6        # to
7        # the ith test point.
8        closest_y = []
9        closest_y = np.argsort(dists[i,:])[:k]
10       t = np.zeros(labelnum,dtype=int)
11       for j in closest_y:
12           t[self.y_train[j]]+=1
13       y_pred[i] = t.argmax()

```

Listing 7: KNN 预测

3.2 SVM

计算 SVM loss、dW(naive):

```
1 dW = np.zeros(W.shape) # initialize the gradient as zero
2
3 # compute the loss and the gradient
4 num_classes = W.shape[1]
5 num_train = X.shape[0]
6 loss = 0.0
7 for i in range(num_train):
8     scores = X[i].dot(W)
9     correct_class_score = scores[y[i]]
10    for j in range(num_classes):
11        if j == y[i]:
12            continue
13        margin = scores[j] - correct_class_score + 1 # note delta = 1
14        if margin > 0:
15            loss += margin
16    loss /= num_train
17
18 # Add regularization to the loss.
19 loss += reg * np.sum(W * W)
20
21 for i in range(num_train):
22     scores = X[i].dot(W)
23     correct_class_score = scores[y[i]]
24     for j in range(num_classes):
25         if j == y[i]:
26             continue
27         margin = scores[j] - correct_class_score + 1 # note delta = 1
28         if margin > 0:
29             dW[:,j] += X[i]
30             dW[:,y[i]] -= X[i]
31 dW /= num_train
32 dW += reg * W * 2
```

Listing 8: SVM(naive)

计算 SVM loss、dW(vectorized):

```
1 loss = 0.0
2 dW = np.zeros(W.shape) # initialize the gradient as zero
3 N = X.shape[0]
4 scores = X.dot(W)
5 score_yi = scores[range(N),y].reshape(-1,1)
6 t = scores - score_yi + 1
7 t[range(N),y] = 0
8 condition = (t>0).astype(int)
9 t = condition*t
10 t = t.sum() / N
11 loss = t + 2 * reg * np.sum(W * W)
```



```

12
13 condition[range(N), y] = - np.sum(condition, axis = 1)
14 dW += np.dot(X.T, condition)/N + 2 * reg * W

```

Listing 9: SVM(vectorized)

训练线性分类器 (batch):

```

1 num_train, dim = X.shape
2 num_classes = (
3     np.max(y) + 1
4 ) # assume y takes values 0...K-1 where K is number of classes
5 if self.W is None:
6     # lazily initialize W
7     self.W = 0.001 * np.random.randn(dim, num_classes)
8
9 # Run stochastic gradient descent to optimize W
10 loss_history = []
11 for it in range(num_iters):
12     X_batch = None
13     y_batch = None
14     indices = np.random.choice(num_train, batch_size)
15     X_batch = X[indices]
16     y_batch = y[indices]
17
18     # evaluate loss and gradient
19     loss, grad = self.loss(X_batch, y_batch, reg)
20     loss_history.append(loss)
21
22     # perform parameter update
23     self.W -= learning_rate*grad

```

Listing 10: 训练线性分类器

SVM 预测:

```

1 y_pred = np.argmax(np.dot(X, self.W), axis = 1)

```

Listing 11: SVM 预测

SVM grid search:

```

1 for i in learning_rates:
2     for j in regularization_strengths:
3         svm = LinearSVM()
4         svm.train(X_train, y_train, learning_rate=i, reg=j,
5                 num_iters=1500, verbose=True)
6         y_train_pred = svm.predict(X_train)
7         y_val_pred = svm.predict(X_val)
8         y_train_acc = np.mean(y_train == y_train_pred)
9         y_val_acc = np.mean(y_val == y_val_pred)
10        results[(i,j)] = (y_train_acc, y_val_acc)
11        if y_val_acc > best_val:

```

```

12         best_val = y_val_acc
13         best_svm = svm

```

Listing 12: SVM grid search

3.3 softmax

Softmax loss, dW(Naive):

```

1  loss = 0.0
2  dW = np.zeros_like(W)
3
4  N , D = X.shape
5  C = W.shape[1]
6
7  for i in range(N):
8      f = np.dot(X[i],W)
9      f -= np.max(f) # avoid overflow
10     loss = loss + np.log(np.sum(np.exp(f))) - f[y[i]]
11     dW[:, y[i]] -= X[i]
12     s = np.exp(f).sum()
13     for j in range(C):
14         dW[:, j] += np.exp(f[j]) / s * X[i]
15 loss = loss / N + reg * np.sum(W * W)
16 dW = dW / N + 2*reg * W

```

Listing 13: Softmax(Naive)

Softmax loss, dW(vectorized):

```

1  N , D = X.shape
2  C = W.shape[1]
3  score = np.dot(X,W)
4  t = np.max(score,axis = 1).reshape(N,1)
5  score -= t
6  s = np.exp(score).sum(axis = 1)
7  loss = - score[range(N),y].sum() + np.log(s).sum()
8  counts = np.exp(score) / s.reshape(N, 1)
9  counts[range(N),y] -= 1
10 dW = np.dot(X.T,counts)
11
12 loss = loss/N + reg * np.sum(W * W)
13
14 dW = dW/N + reg *W

```

Listing 14: Softmax(vectorized)

$$Loss = -\log\left(\frac{e^{s_{y_i}}}{\sum_j e^{s_j}}\right) = -s_{y_i} + \log \sum_j e^{s_j}$$

$$\frac{dL}{dW} = -\frac{ds_{y_i}}{dW} + \frac{1}{\sum_j e^{s_j}} \cdot \left(\sum_j (e^{s_j} \frac{dS_j}{dW})\right)$$

softmax grid search:

```
1 from cs231n.classifiers.linear_classifier import Softmax
2
3 for lr in learning_rates:
4     for rs in regularization_strengths:
5         softmax = Softmax()
6         softmax.train(X_train, y_train, learning_rate = lr, reg=rs,
7             num_iters = 1500,
8                 verbose = True)
9
10        y_pred_train = softmax.predict(X_train)
11        acc_train = np.mean(y_pred_train == y_train)
12
13        y_pred_val = softmax.predict(X_val)
14        acc_val = np.mean(y_pred_val == y_val)
15        results[(lr, rs)] = (acc_train, acc_val)
16
17        if acc_val > best_val:
18            best_val = acc_val
19            best_softmax = softmax
```

Listing 15: softmax grid search

3.4 FC Net

FC 单层向后传播:

```
1 N = x.shape[0]
2 D = x[0].reshape(-1,1).shape[0]
3 M = b.shape[0]
4 # print(N,D,M)
5 out = np.zeros((N,M))
6 for i in range(N):
7     t = x[i].reshape(-1,1)
8     # print(w.shape,t.shape)
9     t = np.dot(w.T,t) + b.reshape(-1,1)
10    out[i,:] = t.T
```

Listing 16: FC 单层前向传播

FC 单层向后传播:

```
1 N = x.shape[0]
2 t = x.shape
3 D = x[0].reshape(-1,1).shape[0]
4 M = b.shape[0]
5
6 x = x.reshape(N,D)
7 dx = np.dot(dout,w.T).reshape(t)
8 dw = np.dot(x.T,dout)
```

```

9      db = dout.sum(axis=0)

```

Listing 17: FC 单层向后传播

Relu 向前传播:

```

1      condition = (x>0).astype(int)
2      out = condition * x

```

Listing 18: Relu 向前传播

Relu 向后传播:

```

1      dx = (x > 0 ).astype(int) * dout

```

Listing 19: Relu 向后传播

SVM loss、dx

```

1      N,C = x.shape
2      scores = x
3      score_yi = scores[range(N),y].reshape(-1,1)
4
5      t = scores - score_yi + 1
6      t[range(N),y] = 0
7      condition = (t>0).astype(int)
8      t = condition*t
9      t = t.sum() / N
10     loss = t
11
12     condition[range(N), y] = - np.sum(condition, axis = 1)
13     # print(condition.shape)
14     dx = condition/N

```

Listing 20: SVM loss、dx

Softmax loss、dx

```

1      N,C = x.shape
2      score = x - np.max(x,axis = 1).reshape(N,1)
3      s = np.exp(score).sum(axis = 1).reshape(N,1)
4      score_yi = score[range(N),y].reshape(N,1)
5      loss = (-score_yi + np.log(s)).sum() /N
6
7
8      expscore = np.exp(score)
9
10     dx = (expscore / s).reshape(N,C)
11
12     dx[range(N),y] -=1
13
14     dx /= N

```

Listing 21: Softmax loss、dx

Two Layer Net Loss:

```
1  scores = None
2  N = X.shape[0]
3  D = X.reshape(N,-1).shape[1]
4  C = self.C
5
6  X_new = X.reshape(N,D)
7
8  # layer 1
9  cp = X.reshape(N,D)
10 cp = np.dot(cp,self.params['W1'])
11 cp = cp + self.params['b1']
12 h = cp
13
14 # relu
15 condition1 = (cp>0).astype(int)
16 cp = condition1 * cp
17 h1 = cp
18 # layer 2
19 cp = np.dot(cp,self.params['W2'])
20 cp = cp + self.params['b2']
21 # softmax
22 scores = cp
23
24
25
26 if y is None:
27     return scores
28
29 loss, grads = 0, {}
30
31 scoresmax = np.max(scores,axis = 1).reshape(N,1)
32 scores = scores - scoresmax
33 expscores = np.exp(scores)
34 t = expscores.sum(axis =1).reshape(N,1)
35 expscores = expscores / t
36 scores_yi = expscores[range(0,N),y]
37 loss = -np.log(scores_yi).sum() /N+ 0.5*self.reg * (self.params['W1'] *
38 self.params['W1']).sum() + 0.5*self.reg * (self.params['W2'] *self.
39 params['W2']).sum()
40
41 # print(loss)
42
43 '''gradient'''
44 # Loss
45 dscore = (expscores).reshape(N,C)
46 dscore[range(N),y] -= 1
47 dscore /= N
48
49 # f2
```

```

48     dw2 = np.dot(h1.T , dscore)
49     db2 = np.sum(dscore,axis = 0)
50     dh1 = np.dot(dscore,self.params['W2'].T)
51
52     # relu
53
54     dh = (h>0).astype(int)
55     dh = dh * dh1
56
57     # f1
58
59     dw1 = np.dot(X_new.T,dh)
60     db1 = np.sum(dh,axis = 0)
61
62     #         dw1 = dw1.reshape(self.params['W1'].shape)
63     #         dw2 = dw2.reshape(self.params['W2'].shape)
64
65     dw1 += self.reg * self.params['W1']
66     dw2 += self.reg * self.params['W2']
67
68     grads['W1']=dw1
69     grads['b1']=db1
70     grads['W2']=dw2
71     grads['b2']=db2

```

Listing 22: Two Layer Net Loss

sgd

```

1 w -= config['learning_rate'] * dw

```

Listing 23: sgd

FC Net grid search:

```

1 input_size = 32 * 32 * 3
2 num_classes = 10
3 best_acc = -1
4 for bs in [200, 400]:
5     for lr in [1e-3, 1e-4, 1e-5]:
6         for hidden_size in [50, 100, 200]:
7             net = TwoLayerNet(input_size, hidden_size, num_classes)
8             solver = Solver(net, data,
9                             num_train_samples=100,
10                            lr_decay=0.9,
11                            num_epochs=20,
12                            print_every=50000,
13                            batch_size = bs,
14                            optim_config={
15                                'learning_rate': lr,
16                            })
17

```

```

18         solver.train()
19
20         # Predict on the validation set
21         val_acc = solver.check_accuracy(data['X_val'], data['y_val'])
22         print ('batch_size = %d, lr = %f, hidden size = %f,
Valid_accuracy: %f' %(bs, lr, hidden_size, val_acc))
23         if val_acc > best_acc:
24             best_acc = val_acc
25             best_model = net

```

Listing 24: FC Net grid search

4 GAT

4.1 theory

对于一个图, 我们将其表示为邻接矩阵 $adj \in R^{N \times N}$, 每个节点 i 都有其特征表示向量 $h_i \in R^D$, 在单个 attention 中, 设 $W \in R^{in \times out}$ 为共有的可学习的参数, 对于直接连接在节点向量上的 attention 层中 $in = D$, $\alpha \in R^{2 \cdot out}$ 为注意力系数向量也是可学习的, 在前向计算中, 隐藏层的向量由以下公式计算:

$$\alpha_{ij} = \frac{\exp \left(\text{LeakyReLU} \left(\vec{a}^T \left[\mathbf{W} \vec{h}_i \| \mathbf{W} \vec{h}_j \right] \right) \right)}{\sum_{k \in \mathcal{N}_i} \exp \left(\text{LeakyReLU} \left(\vec{a}^T \left[\mathbf{W} \vec{h}_i \| \mathbf{W} \vec{h}_k \right] \right) \right)}$$

$$\vec{h}'_i = \sigma \left(\sum_{j \in \mathcal{N}_i} \alpha_{ij} \mathbf{W} \vec{h}_j \right)$$

对于多头 attention, 我们只需设立多个 $W \alpha$, 将计算出的每个节点的隐藏层向量拼接起来即可。

对于此多头 attention 层是最后一层, 就不对其做拼接操作, 而是直接对 K (头数) 个隐藏层向量齐求均值

4.2 codes

GAT 的搭建 (based torch):

```

1  class GAT(nn.Module):
2  def __init__(self, nfeat, nhid, nclass, dropout, alpha, nheads):
3      """Dense version of GAT."""
4      super(GAT, self).__init__()
5      self.dropout = dropout
6
7      self.attentions = [GraphAttentionLayer(nfeat, nhid, dropout=dropout,
alpha=alpha, concat=True) for _ in range(nheads)]

```

```

8         for i, attention in enumerate(self.attentions):
9             self.add_module('attention_{}'.format(i), attention)
10
11         self.out_att = GraphAttentionLayer(nhid * nheads, nclass, dropout=
dropout, alpha=alpha, concat=False) # 第二层(最后一层)的attention
layer
12
13     def forward(self, x, adj):
14         x = F.dropout(x, self.dropout, training=self.training)
15         x = torch.cat([att(x, adj) for att in self.attentions], dim=1) #
cat hidden vector of every node
16         x = F.dropout(x, self.dropout, training=self.training)
17         x = F.elu(self.out_att(x, adj)) # final attention
18         return F.log_softmax(x, dim=1)

```

Listing 25: GAT

```

1     class GraphAttentionLayer(nn.Module):
2         """
3         Simple GAT layer, similar to https://arxiv.org/abs/1710.10903
4         """
5         def __init__(self, in_features, out_features, dropout, alpha, concat=
True):
6             super(GraphAttentionLayer, self).__init__()
7             self.dropout = dropout
8             self.in_features = in_features
9             self.out_features = out_features
10            self.alpha = alpha
11            self.concat = concat
12
13            self.W = nn.Parameter(torch.empty(size=(in_features, out_features)))
14            nn.init.xavier_uniform_(self.W.data, gain=1.414)
15            self.a = nn.Parameter(torch.empty(size=(2*out_features, 1))) #
concat(V, NeigV)
16            nn.init.xavier_uniform_(self.a.data, gain=1.414)
17
18            self.leakyrelu = nn.LeakyReLU(self.alpha)
19
20        def forward(self, h, adj):
21            Wh = torch.mm(h, self.W) # h.shape: (N, in_features), Wh.shape: (N,
out_features)
22            a_input = self._prepare_attentional_mechanism_input(Wh) # 每一个节
点和所有节点, 特征。(Vall, Vall, feature)
23            e = self.leakyrelu(torch.matmul(a_input, self.a).squeeze(2))
24            # 之前计算的是一个节点和所有节点的attention, 实际需要的是连接的节点
的attention系数
25            zero_vec = -9e15*torch.ones_like(e) #
26            attention = torch.where(adj > 0, e, zero_vec) # 无边相连即为-\inf
, 有边相连, 设置值为\alpha_{i,j}将邻接矩阵中小于0的变成负无穷 e~-\inf =
0
27            attention = F.softmax(attention, dim=1) # 按行求softmax。 将系数归

```



```

    _, sum(axis=1) == 1
28     attention = F.dropout(attention, self.dropout, training=self.
        training)
29     h_prime = torch.matmul(attention, Wh)    # 聚合邻居函数, 加权平均
30
31     if self.concat:
32         return F.elu(h_prime)
33     else:
34         return h_prime
35
36 def _prepare_attentional_mechanism_input(self, Wh): # 计算Wh_i || Wh_j
37     N = Wh.size()[0]
38     Wh_repeated_in_chunks = Wh.repeat_interleave(N, dim=0)
39     Wh_repeated_alternating = Wh.repeat(N, 1)
40     all_combinations_matrix = torch.cat([Wh_repeated_in_chunks,
        Wh_repeated_alternating], dim=1)
41     return all_combinations_matrix.view(N, N, 2 * self.out_features)

```

Listing 26: Attention layer

对于训练过程, 我们可以传入训练集的图结构和节点向量矩阵, 在测试的时候将邻接矩阵变幻, 通过 mask 获得测试节点的 label。

```

1     model.train()
2     optimizer.zero_grad()
3
4     features1 = features[0:1000]
5     adj1 = adj[0:1000]
6     adj1 = adj1.T
7     adj1 = adj1[0:1000]
8     adj1 = adj1.T
9     labels1 = labels[0:1000]
10    # print("feature1:", features1.shape)
11    # print("adj1:", adj1.shape)
12    output = model(features1, adj1) # GAT模块
13    # loss_train = F.nll_loss(output[idx_train], labels[idx_train])
14    # acc_train = accuracy(output[idx_train], labels[idx_train])
15
16    loss_train = F.nll_loss(output, labels1)
17    acc_train = accuracy(output, labels1)
18
19
20    # test
21    model.eval()
22    output = model(features, adj)
23    loss_test = F.nll_loss(output[idx_test], labels[idx_test])
24    acc_test = accuracy(output[idx_test], labels[idx_test])

```

Listing 27: train & test